

# Statistical Machine Learning

4주차

담당: 11기 명재성

# Review

---

- Logistic Regression Optimization

$$Loss[Y, \hat{Y}] = - \sum_{i=1}^n [y_i(\boldsymbol{\beta}^T \mathbf{x}_i) - \log(1 + \exp(\boldsymbol{\beta}^T \mathbf{x}_i))]$$

$$\hat{\boldsymbol{\beta}} = \underset{\boldsymbol{\beta}}{argmin} L[Y, \hat{Y}]$$

⇒ Do not have an explicit solution!

# Review

---

- Quadratic approximation (2<sup>nd</sup> order Taylor Expansion)

$$L(\boldsymbol{\theta}) \approx L(\boldsymbol{\theta}_0) + \nabla L(\boldsymbol{\theta}_0)^T (\boldsymbol{\theta} - \boldsymbol{\theta}_0) + \frac{1}{2} (\boldsymbol{\theta} - \boldsymbol{\theta}_0)^T \mathbf{H}(\boldsymbol{\theta}_0) (\boldsymbol{\theta} - \boldsymbol{\theta}_0)$$

where

$$\nabla L(\boldsymbol{\theta}_0) = \left. \frac{\partial}{\partial \boldsymbol{\theta}} L(\boldsymbol{\theta}) \right|_{\boldsymbol{\theta}=\boldsymbol{\theta}_0}$$

$$\mathbf{H}(\boldsymbol{\theta}_0) = \left. \frac{\partial^2}{\partial \boldsymbol{\theta} \partial \boldsymbol{\theta}^T} L(\boldsymbol{\theta}) \right|_{\boldsymbol{\theta}=\boldsymbol{\theta}_0}$$

# Review

---

- Newton-Raphson Method

$$\begin{aligned}\boldsymbol{\theta}^{(t+1)} &= \boldsymbol{\theta}^{(t)} - \mathbf{H}^{-1}(\boldsymbol{\theta}^{(t)}) \nabla L(\boldsymbol{\theta}^{(t)}) \\ &= \boldsymbol{\theta}^{(t)} - \mathbf{H}^{-1}(\boldsymbol{\theta}^{(t)}) \frac{\partial}{\partial \boldsymbol{\theta}^{(t)}} L(\boldsymbol{\theta}^{(t)})\end{aligned}$$

$$cf. \quad \theta^{(t+1)} = \theta^{(t)} - \frac{f'(\theta^{(t)})}{f''(\theta^{(t)})}$$

# Review

---

$$L[\boldsymbol{\beta}] = - \sum_{i=1}^n [y_i(\boldsymbol{\beta}^T \mathbf{X}_i) - \log(1 + \exp(\boldsymbol{\beta}^T \mathbf{X}_i))]$$

$$\nabla L(\boldsymbol{\beta}) = \frac{\partial}{\partial \boldsymbol{\beta}} L(\boldsymbol{\beta}) = - \sum_{i=1}^n \left[ y_i \mathbf{X}_i - \frac{\exp(\boldsymbol{\beta}^T \mathbf{X}_i)}{1 + \exp(\boldsymbol{\beta}^T \mathbf{X}_i)} \mathbf{X}_i \right]$$

$$\mathbf{H}(\boldsymbol{\beta}) = \frac{\partial^2}{\partial \boldsymbol{\beta} \partial \boldsymbol{\beta}^T} L(\boldsymbol{\beta}) = \sum_{i=1}^n \left[ \left( \frac{\exp(\boldsymbol{\beta}^T \mathbf{X}_i)}{1 + \exp(\boldsymbol{\beta}^T \mathbf{X}_i)} \right) \left( \frac{1}{1 + \exp(\boldsymbol{\beta}^T \mathbf{X}_i)} \right) \mathbf{X}_i \mathbf{X}_i^T \right]$$

# Review

---

$$L[\boldsymbol{\beta}] = - \sum_{i=1}^n [y_i(\boldsymbol{\beta}^T \mathbf{x}_i) - \log(1 + \exp(\boldsymbol{\beta}^T \mathbf{x}_i))]$$

Update

$$\boldsymbol{\beta}^{(t+1)} = \boldsymbol{\beta}^{(t)} - \mathbf{H}^{-1}(\boldsymbol{\beta}^{(t)}) \nabla L(\boldsymbol{\beta}^{(t)})$$

until

$$\|\boldsymbol{\beta}^{t+1} - \boldsymbol{\beta}^t\| < \epsilon \quad \text{for small } \epsilon > 0$$

# Review

| Solvers                      |             |         |             |       |        |
|------------------------------|-------------|---------|-------------|-------|--------|
| Penalties                    | 'liblinear' | 'lbfgs' | 'newton-cg' | 'sag' | 'saga' |
| Multinomial + L2 penalty     | no          | yes     | yes         | yes   | yes    |
| OVR + L2 penalty             | yes         | yes     | yes         | yes   | yes    |
| Multinomial + L1 penalty     | no          | no      | no          | no    | yes    |
| OVR + L1 penalty             | yes         | no      | no          | no    | yes    |
| Behaviors                    |             |         |             |       |        |
| Penalize the intercept (bad) | yes         | no      | no          | no    | no     |
| Faster for large datasets    | no          | no      | no          | yes   | yes    |
| Robust to unscaled datasets  | yes         | yes     | yes         | no    | no     |

# Review

---

- Gradient Descent

$$\boldsymbol{\theta}^{(t+1)} = \boldsymbol{\theta}^{(t)} - \mathbf{H}^{-1}(\boldsymbol{\theta}^{(t)}) \nabla L(\boldsymbol{\theta}^{(t)})$$

$$\Rightarrow \boldsymbol{\theta}^{(t+1)} = \boldsymbol{\theta}^{(t)} - \eta^{(t)} \nabla L(\boldsymbol{\theta}^{(t)})$$



# Review

---

- Gradient Descent

$$\boldsymbol{\theta}^{(t+1)} = \boldsymbol{\theta}^{(t)} - \mathbf{H}^{-1}(\boldsymbol{\theta}^{(t)}) \nabla L(\boldsymbol{\theta}^{(t)})$$


$$\Rightarrow \boldsymbol{\theta}^{(t+1)} = \boldsymbol{\theta}^{(t)} - \eta^{(t)} \nabla L(\boldsymbol{\theta}^{(t)}) \quad \rightarrow \text{ in Deep Learning}$$

$$\Rightarrow \boldsymbol{\theta}^{(t+1)} = \boldsymbol{\theta}^{(t)} - \eta \nabla L(\boldsymbol{\theta}^{(t)})$$

# Review

---

- (Batch) Gradient Descent




- Batch size =  $n$

# Review

---

- Stochastic Gradient Descent




- Batch size = 1

# Review


---

- Mini-Batch Gradient Descent



- Batch size =  $p$  ,      where  $p \times k = n$

# Review



# Stein's Paradox

---

- Let  $\mathbf{X} = [X_1, \dots, X_p]^T \sim N_p(\boldsymbol{\theta}, I)$
- The UMVUE and MLE of  $\boldsymbol{\theta}$  is

$$\hat{\boldsymbol{\theta}}_{MLE,UMVUE} = \mathbf{X}$$

- Using squared error loss, the risk of  $\hat{\boldsymbol{\theta}}_{MLE,UMVUE}$  is

$$R(\boldsymbol{\theta}, \hat{\boldsymbol{\theta}}_{UMVUE}) = E[||\mathbf{X} - \boldsymbol{\theta}||^2] = p$$

# Stein's Paradox

---

- James and Stein (1961) Estimator

$$\hat{\boldsymbol{\theta}}_{JS} = \left(1 - \frac{p-2}{\|\mathbf{X}\|^2}\right) \mathbf{X}$$

- When  $p \geq 3$ ,

$$R(\boldsymbol{\theta}, \hat{\boldsymbol{\theta}}_{JS}) = p - (p-2)E\left(\frac{1}{\|\mathbf{X}\|^2}\right) < p$$

# Steins Paradox

---

- Proof

$$\begin{aligned} R(\boldsymbol{\theta}, \hat{\boldsymbol{\theta}}_{JS}) &= E \left[ \left\| \mathbf{X} - \boldsymbol{\theta} - \frac{(p-2)\mathbf{X}}{\|\mathbf{X}\|^2} \right\|^2 \right] \\ &= p - 2(p-2) \sum_j^p E \left( \frac{X_j(X_j - \theta_j)}{\|\mathbf{X}\|^2} \right) + (p-2)^2 E \left( \frac{1}{\|\mathbf{X}\|^2} \right) \\ &= p - (p-2) E \left( \frac{1}{\|\mathbf{X}\|^2} \right) \end{aligned}$$

Since  $\sum_j^p E \left( \frac{X_j(X_j - \theta_j)}{\|\mathbf{X}\|^2} \right) = (p-2) E \left( \frac{1}{\|\mathbf{X}\|^2} \right)$



# Steins Paradox

---

- JS estimator shrinks each component of  $\mathbf{X}$  towards the origin, and thus the biggest improvement comes when  $||\boldsymbol{\theta}||$  is close to zero.
- Normality assumption is not critical, and similar results can be shown for a wide class of distributions.

# Ridge Regression

---

- Normal Equation

$$(\mathbf{X}^T \mathbf{X}) \boldsymbol{\beta} = \mathbf{X}^T \mathbf{Y}$$

- The OLS estimator

$$\hat{\boldsymbol{\beta}}_{OLS} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$$

# Ridge Regression

---

- Normal Equation

$$(\mathbf{X}^T \mathbf{X}) \boldsymbol{\beta} = \mathbf{X}^T \mathbf{Y}$$

- The OLS estimator

$$\hat{\boldsymbol{\beta}}_{OLS} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$$

- What if  $\mathbf{X}^T \mathbf{X}$  is not invertible?

# Ridge Regression

---

- We can consider

$$\hat{\boldsymbol{\beta}}_{Ridge} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{Y}$$

- Ridge estimator is

$$\hat{\boldsymbol{\beta}}_{Ridge} = \underset{\boldsymbol{\beta}}{\operatorname{argmin}} (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta}) + \lambda \boldsymbol{\beta}^T \boldsymbol{\beta}$$

# Lagrange Multiplier Theorem

---

- Primal Problem

$$\min_{\boldsymbol{\beta}} (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta})$$

$$\text{subject to } \|\boldsymbol{\beta}\|_2^2 \leq C, \quad \text{where } \|\boldsymbol{\beta}\|_2^2 = \boldsymbol{\beta}^T \boldsymbol{\beta} = \sum_j^p \beta_j^2$$

- Dual Problem

$$\min_{\boldsymbol{\beta}} (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta}) + \lambda (\|\boldsymbol{\beta}\|_2^2 - C)$$

# Lagrange Multiplier Theorem


---

$$(\hat{\beta}^{\lambda,2} =) \hat{\beta}_{Ridge} = \underset{\beta}{\operatorname{argmin}} (\mathbf{Y} - \mathbf{X}\beta)^T (\mathbf{Y} - \mathbf{X}\beta) + \lambda \|\beta\|_2^2$$


$$\Leftrightarrow \underset{\beta}{\operatorname{argmin}} (\mathbf{Y} - \mathbf{X}\beta)^T (\mathbf{Y} - \mathbf{X}\beta) + \lambda (\|\beta\|_2^2 - C)$$

```
# Logistic regression
from sklearn.linear_model import LogisticRegression
Logit = LogisticRegression(C=1e2, random_state=1023) # C = 1/λ. 디폴트: L2, One-versus-Rest.
Logit.fit(X_train_std, y_train)
```

# Ridge Regression



# Ridge Regression





# Bias-Variance Trade off

---

- Expected Prediction Error

$$E[(Y_0 - \hat{Y}_0)^2] = \sigma^2 + E[(\mu_0 - \hat{Y}_0)^2]$$

Irreducible error

model error

where  $Y_0 = \mu_0 + \epsilon_0 = \mathbf{x}_0^T \boldsymbol{\beta} + \epsilon_0$

and  $\hat{Y}_0 = \mathbf{x}_0^T \hat{\boldsymbol{\beta}}$

# Bias-Variance Trade off

---

- Model Error

$$\begin{aligned} E[(\mu_0 - \hat{Y}_0)^2] &= E[(\mu_0 - E[\hat{Y}_0] + E[\hat{Y}_0] - \hat{Y}_0)^2] \\ &= \underbrace{(\mu_0 - E[\hat{Y}_0])^2}_{\text{Bias}^2} + \underbrace{\text{Var}[\hat{Y}_0]}_{\text{variance}} \end{aligned}$$

# LASSO Regression

---

- Ridge Regression solves

$$\min_{\beta} (\mathbf{Y} - \mathbf{X}\beta)^T (\mathbf{Y} - \mathbf{X}\beta) + \lambda \|\beta\|_2^2 \quad (L2 \text{ penalty})$$

- LASSO Regression solves

# LASSO Regression

---

- Ridge Regression solves

$$\min_{\boldsymbol{\beta}} (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta}) + \lambda ||\boldsymbol{\beta}||_2^2 \quad (L2 \text{ penalty})$$

- LASSO Regression solves

$$\min_{\boldsymbol{\beta}} (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta}) + \lambda ||\boldsymbol{\beta}||_1 \quad (L1 \text{ penalty})$$

# LASSO Regression


---

- LASSO (Least Absolute Shrinkage and Selection Operator)


$$(\hat{\boldsymbol{\beta}}^{\lambda,1} =) \hat{\boldsymbol{\beta}}_{LASSO} = \underset{\boldsymbol{\beta}}{argmin} (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta}) + \lambda ||\boldsymbol{\beta}||_1$$

$$\text{where } ||\boldsymbol{\beta}||_1 = \sum_j^p |\beta_j|$$

# LASSO Regression



# LASSO Regression



# One-dimensional Case

---

- For simplicity, let  $y_i = \beta x_i + \epsilon_i$ , where  $\sum_i^n x_i = 0$  and  $\sum_i^n x_i^2 = n$
- Least Square estimator is

$$\hat{\beta}_{OLS} = \frac{1}{n} \sum x_i y_i$$

- Ridge estimator is

$$\hat{\beta}_{Ridge} = \frac{\hat{\beta}_{OLS}}{1 + \lambda}$$




# One-dimensional Case

## Definition (Subgradient)

For a convex function  $f : \mathbb{R}^p \rightarrow \mathbb{R}$ , a vector  $\mathbf{z} \in \mathbb{R}^p$  is to be a *subgradient* of  $f$  at  $\beta$  if

$$f(\beta') \geq f(\beta) + \mathbf{z}^T (\beta' - \beta) \quad \text{for all } \beta' \in \mathbb{R}^p.$$



# One-dimensional Case

---

- For LASSO, the objective function is

$$f(\beta) = \frac{1}{n} \sum (y_i - \beta x_i)^2 + \lambda |\beta|$$

and its subdifferential is

$$\partial f(\beta) = \begin{cases} \beta - \hat{\beta}_{OLS} + \lambda & \text{if } \beta > 0 \\ \beta - \hat{\beta}_{OLS} + \lambda[-1, 1] & \text{if } \beta = 0 \\ \beta - \hat{\beta}_{OLS} - \lambda & \text{if } \beta < 0 \end{cases}$$

# One-dimensional Case

---

- LASSO estimator is

$$\hat{\beta}_{LASSO} = \begin{cases} \hat{\beta}_{OLS} - \lambda & \text{if } \hat{\beta}_{OLS} > \lambda \\ 0 & \text{if } |\hat{\beta}_{OLS}| \leq \lambda \\ \hat{\beta}_{OLS} + \lambda & \text{if } \hat{\beta}_{OLS} < -\lambda \end{cases}$$

# One-dimensional Case

---

- LASSO estimator is

$$\hat{\beta}_{LASSO} = \begin{cases} \hat{\beta}_{OLS} - \lambda & \text{if } \hat{\beta}_{OLS} > \lambda \\ 0 & \text{if } |\hat{\beta}_{OLS}| \leq \lambda \\ \hat{\beta}_{OLS} + \lambda & \text{if } \hat{\beta}_{OLS} < -\lambda \end{cases}$$

- Soft-thresholding operator

$$S_{\lambda}(x) = \text{sign}(x) (|x| - \lambda)_+$$

# One-dimensional Case

---


- LASSO estimator is

$$\hat{\beta}_{LASSO} = S_{\lambda}(\hat{\beta}_{OLS})$$


- Soft-thresholding operator

$$S_{\lambda}(x) = \text{sign}(x) (|x| - \lambda)_+$$


# One-dimensional Case



(a) Hard Thresh.



(b) Ridge Regression



(c) Lasso (Soft Thresh.)

# Feature Selection and Extraction

---

- LASSO (Least Absolute Shrinkage and Selection Operator)

$$(\hat{\boldsymbol{\beta}}^{\lambda,1} =) \hat{\boldsymbol{\beta}}_{LASSO} = \underset{\boldsymbol{\beta}}{argmin} (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta}) + \lambda ||\boldsymbol{\beta}||_1$$

$$\text{where } ||\boldsymbol{\beta}||_1 = \sum_j^p |\beta_j|$$

# Feature Selection and Extraction

---

- LASSO (Least Absolute Shrinkage and Selection Operator)

$$(\hat{\boldsymbol{\beta}}^{\lambda,1} =) \hat{\boldsymbol{\beta}}_{LASSO} = \underset{\boldsymbol{\beta}}{argmin} (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta}) + \lambda ||\boldsymbol{\beta}||_1$$

$$\text{where } ||\boldsymbol{\beta}||_1 = \sum_j^p |\beta_j|$$



# Feature Selection and Extraction

---

- LASSO (Least Absolute Shrinkage and Selection Operator)
- LASSO estimator gives a **sparse** solution

⇒ Thus, features are selected automatically!

# Feature Selection and Extraction

---

- Feature Selection
  - Subset selection, Stepwise method, LASSO, Least Angle Regression etc..
- Feature Extraction (Dimension Reduction)
  - Principal Component Analysis, Partial Least Square, Discriminant Analysis, Factor Analysis, Latent Class Analysis, etc..

# Elastic Net

---


- Elastic Net solves

$$\min_{\boldsymbol{\beta}} (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{Y} - \mathbf{X}\boldsymbol{\beta}) + \lambda \left[ \alpha \|\boldsymbol{\beta}\|_1 + \frac{1}{2} (1 - \alpha) \|\boldsymbol{\beta}\|_2^2 \right]$$

⇒ middle ground of LASSO and Ridge penalty

# Elastic Net


- $L_p$  penalty



**FIGURE 3.12.** *Contours of constant value of  $\sum_j |\beta_j|^q$  for given values of  $q$ .*

# Elastic Net


- Elastic Net



**FIGURE 3.13.** Contours of constant value of  $\sum_j |\beta_j|^q$  for  $q = 1.2$  (left plot), and the elastic-net penalty  $\sum_j (\alpha \beta_j^2 + (1-\alpha)|\beta_j|)$  for  $\alpha = 0.2$  (right plot). Although visually very similar, the elastic-net has sharp (non-differentiable) corners, while the  $q = 1.2$  penalty does not.


# Elastic Net

- Elastic Net



**Figure 4.2** The elastic-net ball with  $\alpha = 0.7$  (left panel) in  $\mathbb{R}^3$ , compared to the  $\ell_1$  ball (right panel). The curved contours encourage strongly correlated variables to share coefficients (see Exercise 4.2 for details).

# Elastic Net



**Figure 4.1** Six variables, highly correlated in groups of three. The lasso estimates ( $\alpha = 1$ ), as shown in the left panel, exhibit somewhat erratic behavior as the regularization parameter  $\lambda$  is varied. In the right panel, the elastic net with ( $\alpha = 0.3$ ) includes all the variables, and the correlated groups are pulled together.

# Elastic Net

```
from sklearn.linear_model import LogisticRegression

lr2_1 = LogisticRegression(penalty='l2', C=1.0)    # L2 with  $C(=1/\lambda)=1$ 
lr2_0_1 = LogisticRegression(penalty='l2', C=0.1) # L2 with  $C(=1/\lambda)=0.1$ 

lr1_1 = LogisticRegression(penalty='l1', C=1.0)    # L1 with  $C(=1/\lambda)=1$ 
lr1_0_1 = LogisticRegression(penalty='l1', C=0.1) # L1 with  $C(=1/\lambda)=0.1$ 

lre_1 = LogisticRegression(penalty='elasticnet', C=1.0, l1_ratio=0.2) # Elasticnet with  $C(=1/\lambda)=1.0$ 
lre_0_1 = LogisticRegression(penalty='elasticnet', C=0.1, l1_ratio=0.2) # Elasticnet with  $C(=1/\lambda)=0.1$ 
```

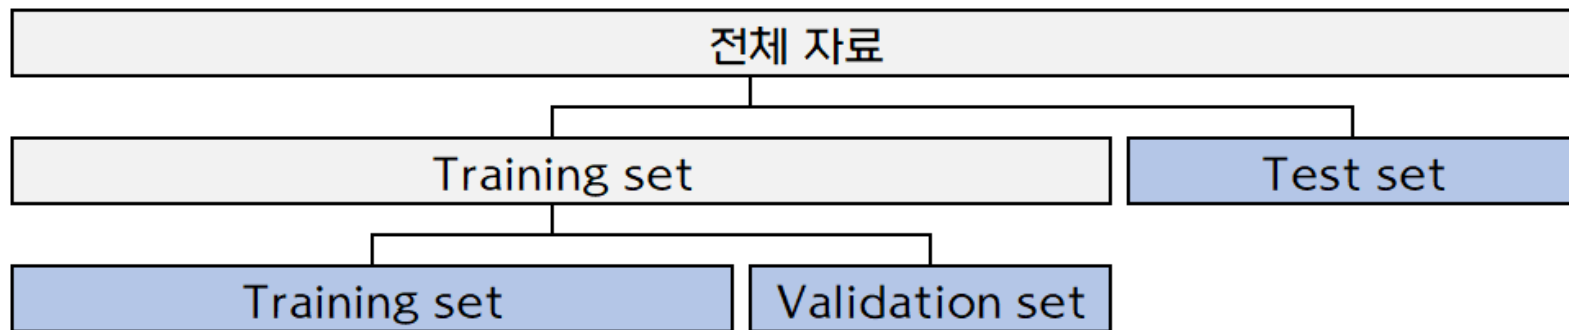
```
library(glmnet)
```

```
fit = glmnet(x, y, alpha = 0.2, weights = c(rep(1,50),rep(2,50)), nlambdas = 20)
```



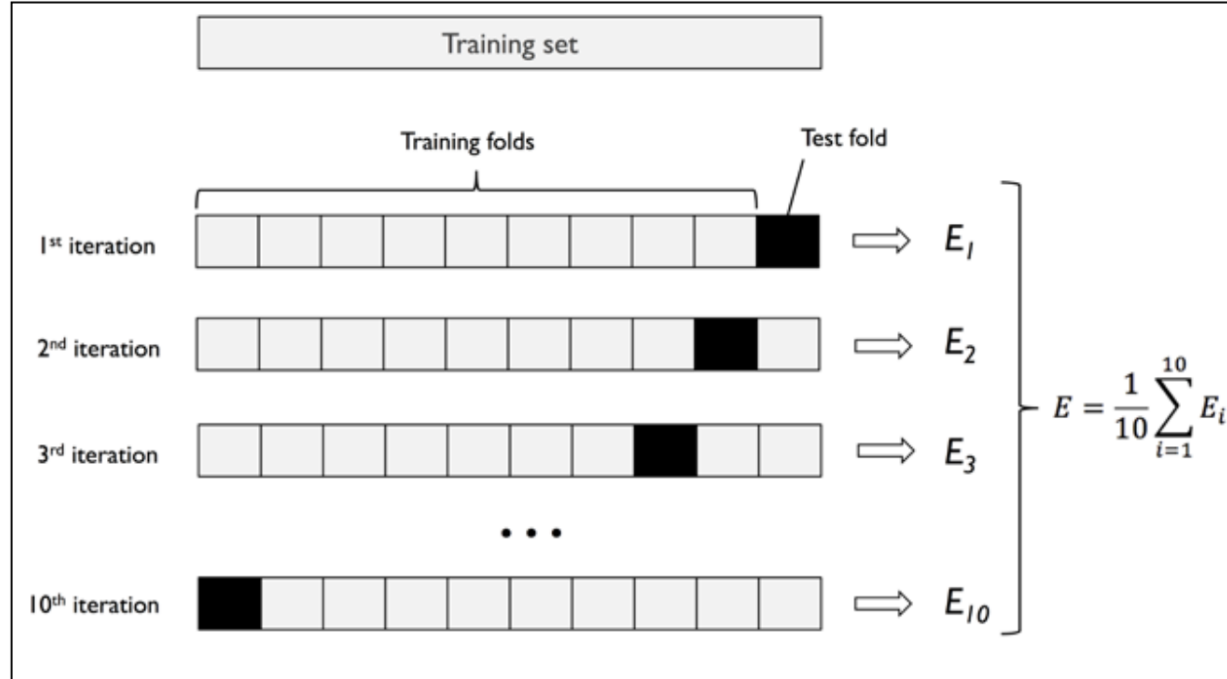
# Cross Validation

---



# K-fold Cross Validation

- $K = 10$



# K-fold Cross Validation

---

```
### Pipeline Streaming: 표준화 → PCA → Logistic Regression ###
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import make_pipeline
pipe_lr = make_pipeline(StandardScaler(),
                        PCA(n_components=4),
                        LogisticRegression(random_state=1, solver='lbfgs')) # 적용 순서대로 나열
pipe_lr.fit(X_train, y_train) # 표준화(fit → transform) → PCA(fit → transform) → Logistic Reg fit의 순서로 처리
```

# K-fold Cross Validation

---


```
[ ] ### K-fold cross-validation using pipeline ###  
from sklearn.model_selection import cross_val_score  
scores = cross_val_score(estimator=pipe_lr, X=X_train, y=y_train, cv=10) # Accuracy scores  
print('CV accuracy scores: %s' % scores)  
  
import numpy as np  
print('CV accuracy: %.3f +/- %.3f' % (np.mean(scores), np.std(scores)))
```

➡ CV accuracy scores: [0.97826087 0.95652174 0.95652174 0.95652174 0.91304348 0.95555556  
0.97777778 0.97777778 1. 0.97777778]  
CV accuracy: 0.965 +/- 0.022

# Nested Cross Validation

- $K_1 = 5$

$$K_2 = 2$$



# Grid Search CV

```
[ ] # Decision tree
    from sklearn.tree import DecisionTreeClassifier
    from sklearn.model_selection import GridSearchCV
    from sklearn.model_selection import KFold
    inner_cv=KFold(n_splits=3, shuffle=True, random_state=0)
    outer_cv=KFold(n_splits=5, shuffle=True, random_state=0)
    gs = GridSearchCV(estimator=DecisionTreeClassifier(random_state=0),
                      param_grid=[{'max_depth': [1, 2, 3, 4, 5, 6, 7, None]}],
                      scoring='accuracy', cv=inner_cv)
    scores = cross_val_score(gs, X, y, scoring='accuracy', cv=outer_cv)
    print('CV accuracy: %.3f +/- %.3f' % (np.mean(scores), np.std(scores)))
```

➡ CV accuracy: 0.942 +/- 0.012

# Grid Search CV

---

`cv.glmnet {glmnet}`

R Documentation

## Cross-validation for glmnet

### Description

Does k-fold cross-validation for glmnet, produces a plot, and returns a value for `lambda` (and `gamma` if `relax=TRUE`)

### Usage

```
cv.glmnet(x, y, weights = NULL, offset = NULL, lambda = NULL,
  type.measure = c("default", "mse", "deviance", "class", "auc", "mae",
    "C"), nfolds = 10, foldid = NULL, alignment = c("lambda",
    "fraction"), grouped = TRUE, keep = FALSE, parallel = FALSE,
  gamma = c(0, 0.25, 0.5, 0.75, 1), relax = FALSE, trace.it = 0, ...)
```

# reference

자료

19-2 STAT424 통계적 머신러닝 - 박유성 교수님

교재

파이썬을 이용한 통계적 머신러닝 (2020) - 박유성

ISLR (2013) - G. James, D. Witten, T. Hastie, R. Tibshirani

The elements of Statistical Learning (2001) - J. Friedman, T. Hastie, R. Tibshirani

Statistical Learning with Sparsity: The Lasso and Generalizations(2015)

- R. Tibshirani, , T. Hastie, M. Wainwright